

# Zeppelin

FINISHED ▶ ⌵ 📖 ⚙️

```
%pyspark
from pandas import Series, DataFrame
import pandas as pd
import numpy as np

frame = DataFrame({'data1': np.random.randn(1000), 'data2': np.random.randn(1000)})
factor = pd.cut(frame.data1,4)
factor[:10]
```

```
0      (-0.0938, 1.475]
1      (-1.663, -0.0938]
2      (1.475, 3.0437]
3      (-0.0938, 1.475]
4      (-0.0938, 1.475]
5      (-1.663, -0.0938]
6      (-0.0938, 1.475]
7      (-3.238, -1.663]
8      (1.475, 3.0437]
9      (-1.663, -0.0938]
Name: data1, dtype: category
Categories (4, object): [(-3.238, -1.663] < (-1.663, -0.0938] < (-0.0938, 1.475] < (1.475, 3.0437]]
```

FINISHED ▶ ⌵ 📖 ⚙️

```
%pyspark
def get_stats(group):
    return {'min': group.min(), 'max': group.max(), 'count': group.count(), 'mean': group.mean()}
grouped = frame.data2.groupby(factor)
grouped.apply(get_stats).unstack()
```

	count	max	mean	min
data1				
(-3.238, -1.663]	51.0	1.847021	-0.016305	-2.203226
(-1.663, -0.0938]	439.0	3.036204	-0.049544	-3.537838
(-0.0938, 1.475]	447.0	2.747356	0.055262	-3.404124
(1.475, 3.0437]	63.0	1.986052	0.005977	-3.107897

FINISHED ▶ ⌵ 📖 ⚙️

```
%pyspark
# compute quantile numbers
grouping = pd.qcut(frame.data1, 10, labels=False)
grouped = frame.data2.groupby(grouping)
grouped.apply(get_stats).unstack()
```

count max mean min

	count	max	mean	min
data1				
0	100.0	2.372078	-0.064840	-3.537838
1	100.0	3.036204	0.060842	-2.300333
2	100.0	2.931313	-0.065915	-2.357223
3	100.0	2.939984	-0.091078	-2.821557
4	100.0	1.954144	-0.034637	-1.940848
5	100.0	2.181723	-0.076362	-2.512935
6	100.0	2.139447	-0.123829	-3.404124
7	100.0	2.366889	0.125820	-1.946764
8	100.0	2.747356	0.159315	-1.940647
9	100.0	2.095359	0.135655	-3.107897

```
%pyspark
from pandas import Series, DataFrame
import pandas as pd
import numpy as np
s = Series(np.random.randn(6))
s[::2] = np.nan
s
```

FINISHED ▶ ⌵ 📖 ⚙

```
0      NaN
1    1.114197
2      NaN
3   -0.402738
4      NaN
5    0.749553
dtype: float64
```

```
%pyspark
s.fillna(s.mean())
```

ERROR ▶ ⌵ 📖 ⚙

```
%pyspark
states = ['Ohio', 'New York', 'Vermont', 'Florida', 'Oregon', 'Nevada', 'California', 'Idaho']
group_key = ['East'] * 4 + ['West'] * 4
data = Series(np.random.randn(8), index=states)
data[['Vermont', 'Nevada', 'Idaho']] = np.nan
data
```

FINISHED ▶ ⌵ 📖 ⚙

```
Ohio      0.742188
New York  0.416759
Vermont    NaN
Florida   -2.267450
Oregon    -0.935032
Nevada     NaN
Califronia 0.168355
Idaho      NaN
dtype: float64
```

```
%pyspark
data.groupby(group_key).mean()
```

FINISHED ▶ ⌵ 📖 ⚙️

```
East      -0.369501
West      -0.383338
dtype: float64
```

```
%pyspark
fill_mean = lambda g : g.fillna(g.mean())
data.groupby(group_key).apply(fill_mean)
```

FINISHED ▶ ⌵ 📖 ⚙️

```
Ohio      0.742188
New York  0.416759
Vermont   -0.369501
Florida   -2.267450
Oregon    -0.935032
Nevada    -0.383338
Califronia 0.168355
Idaho     -0.383338
dtype: float64
```

```
%pyspark
fill_values = {'East': 0.5, 'West': -1}
fill_func = lambda g: g.fillna(fill_values[g.name])
data.groupby(group_key).apply(fill_func)
```

FINISHED ▶ ⌵ 📖 ⚙️

```
Ohio      0.742188
New York  0.416759
Vermont    0.500000
Florida   -2.267450
Oregon    -0.935032
Nevada    -1.000000
Califronia 0.168355
Idaho     -1.000000
dtype: float64
```

```
%pyspark
```

FINISHED ▶ ⌵ 📖 ⚙️

```
from pandas import Series, DataFrame
import pandas as pd
import numpy as np

df = DataFrame({'category': ['a','a','a','a','b','b','b','b'], 'data': np.random.randn(8),
df
```

	category	data	weights
0	a	0.410116	0.378536
1	a	-1.852803	0.336846
2	a	1.033562	0.800144
3	a	1.130160	0.615388
4	b	0.531768	0.985667
5	b	0.388601	0.749693
6	b	-0.515092	0.900023
7	b	-0.528385	0.986466

```
%pyspark
grouped = df.groupby('category')
get_wavg = lambda g: np.average(g['data'], weights=g['weights'])
grouped.apply(get_wavg)
```

FINISHED ▶ ⌵ 📖 ⚙️

```
category
a    0.494445
b   -0.046758
dtype: float64
```



```
%pyspark
close_px = pd.read_csv('/Users/Rhon/Desktop/Capstone/MorganStanley.csv', parse_dates=True,
close_px.info()
```

FINISHED ▶ ⌵ 📖 ⚙️

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 4276 entries, 2017-03-29 to 2000-03-30
Data columns (total 6 columns):
Open          4276 non-null float64
High          4276 non-null float64
Low           4276 non-null float64
Close         4276 non-null float64
Volume        4276 non-null int64
Adj Close     4276 non-null float64
dtypes: float64(5), int64(1)
memory usage: 233.8 KB
```

```
%pyspark
close_px[-4:]
```

FINISHED ▶ ⌵ 📖 ⚙️

	Open	High	Low	Close	Volume	Adj Close	↓
Date							
2000-04-04	86.2500	86.250	74.1250	78.625000	8308000	49.772307	
2000-04-03	83.4375	86.750	83.4375	85.375000	3947200	54.045287	
2000-03-31	83.5000	84.125	79.5000	82.875000	4280200	52.462702	
2000-03-30	85.5625	88.250	82.8125	84.140602	3842200	53.263871	

```
%pyspark
rets = close_px.pct_change().dropna()
spx_corr = lambda x: x.corrwith(x['High'])
by_year = rets.groupby(lambda x: x.year)
by_year.apply(spx_corr)
```

FINISHED ▶ 🔍 📖 ⚙️

	Open	High	Low	Close	Volume	Adj Close	↓
2000	0.669288	1.0	0.690060	0.667640	0.106232	0.668156	
2001	0.630900	1.0	0.745442	0.718930	0.064977	0.718837	
2002	0.668289	1.0	0.713400	0.665421	-0.031480	0.664664	
2003	0.600475	1.0	0.652235	0.652582	0.189158	0.655030	
2004	0.670939	1.0	0.731261	0.649752	0.043832	0.647195	
2005	0.666337	1.0	0.653405	0.620970	0.248077	0.621446	
2006	0.698403	1.0	0.651263	0.622167	0.195949	0.622308	
2007	0.794187	1.0	0.809176	0.732038	0.020661	0.528513	
2008	0.837269	1.0	0.746830	0.657567	-0.085290	0.657144	
2009	0.658649	1.0	0.770845	0.673992	0.228678	0.675355	
2010	0.693147	1.0	0.729389	0.664821	0.131331	0.663785	
2011	0.763674	1.0	0.775016	0.665047	0.113350	0.665513	
2012	0.759505	1.0	0.800502	0.699982	0.204002	0.699265	
2013	0.694435	1.0	0.701882	0.647986	0.267003	0.647766	
2014	0.742903	1.0	0.783281	0.748653	0.165451	0.748578	
2015	0.774731	1.0	0.796929	0.719418	-0.030441	0.719371	
2016	0.751636	1.0	0.825287	0.728162	-0.001302	0.728161	

```
%pyspark

# Annual correlation of Apple with Morgan Stanley
by_year.apply(lambda g: g['Low'].corr(g['Volume']))
```

FINISHED ▶ 🔍 📖 ⚙️

```

2000    -0.263173
2001    -0.220835
2002    -0.262384
2003    -0.070999
2004    -0.309774
2005    -0.167750
2006    -0.228065
2007    -0.291658
2008    -0.322245
2009    -0.097808
2010    -0.209288
2011    -0.198386
2012    -0.069018
2013    -0.104618
2014    -0.163641
2015    -0.294724
2016    -0.215895
2017    -0.271866

```



```
%pyspark
```

FINISHED ▶ ⌵ 📖 ⚙️

```
# applying Ordinary Least Squares (OLS) regression on each chunk of data
```

```

import statsmodels.api as sm
def regression(data, yvar, xvars):
    Y = data[yvar]
    X = data[xvars]
    X['intercept'] = 1.
    result = sm.OLS(Y,X).fit()
    return result.params

```

```
by_year.apply(regression, 'High', ['Open'])
```

```

      Open  intercept
2000  0.562084   0.000418
2001  0.548357   0.000787
2002  0.556115   0.000468
2003  0.530821  -0.000462
2004  0.581484   0.000049
2005  0.574285  -0.000072
2006  0.564277  -0.000543
2007  0.701772   0.000531
2008  0.618472   0.001565
2009  0.569047  -0.000830
2010  0.547569   0.000237
2011  0.677649   0.000792
2012  0.615043  -0.000235
2013  0.639176  -0.000614
2014  0.645339  -0.000292
2015  0.649278   0.000250
2016  0.682128  -0.000211

```



```
%pyspark
import timeit
start = timeit.timeit()
close_px[-4:]
end = timeit.timeit()
print(end - start)
```

0.00579380989075

```
%pyspark
import timeit
start = timeit.timeit()

rets = close_px.pct_change().dropna()
spx_corr = lambda x: x.corrwith(x['High'])
by_year = rets.groupby(lambda x: x.year)
by_year.apply(spx_corr)

end = timeit.timeit()
print(end - start)
```

0.00370502471924

```
%pyspark
import timeit
start = timeit.timeit()

by_year.apply(lambda g: g['Low'].corr(g['Volume']))

end = timeit.timeit()
print(end - start)
```

-0.00384306907654

```
%pyspark
import timeit
start = timeit.timeit()

# applying Ordinary Least Squares (OLS) regression on each chunk of data

import statsmodels.api as sm
def regression(data, yvar, xvars):
    Y = data[yvar]
    X = data[xvars]
    X['intercept'] = 1.
    result = sm.OLS(Y,X).fit()
    return result.params

by_year.apply(regression, 'High', ['Open'])
```

```
end = timeit.timeit()  
print(end - start)
```

-0.00264191627502

READY ▶ ⌵ 📖 ⚙

↓